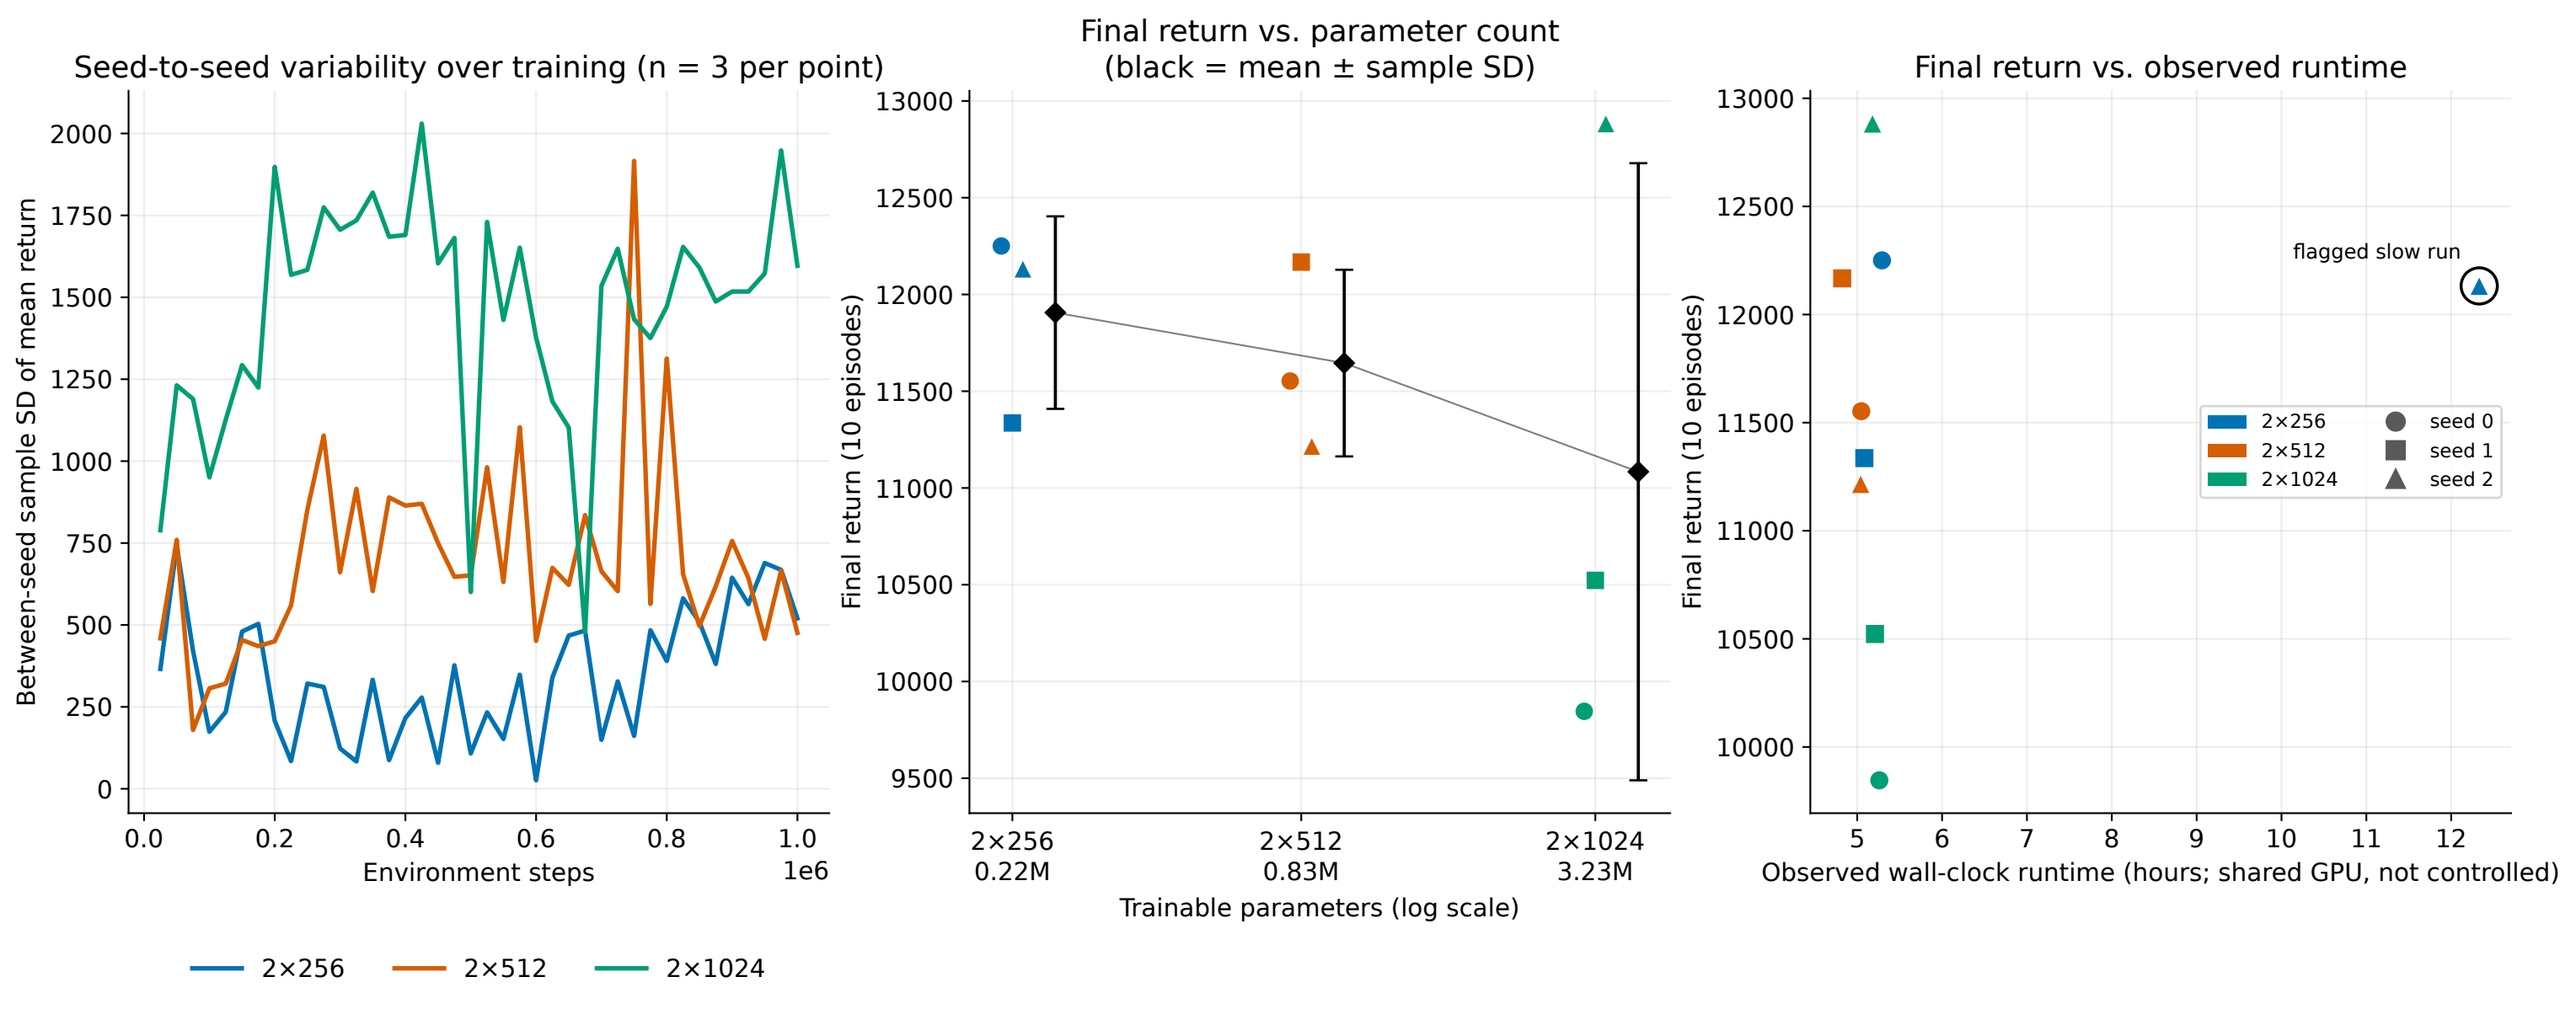
Seed-to-seed variability over training (n = 3 per point)
Between-seed sample SD of mean return
Environment steps
1e6
2×256
2×512
2×1024
Final return vs. parameter count
(black = mean ± sample SD)
Final return (10 episodes)
2×256
0.22M
2×512
0.83M
2×1024
3.23M
Trainable parameters (log scale)
Final return vs. observed runtime
flagged slow run
seed 0
seed 1
seed 2
Observed wall-clock runtime (hours; shared GPU, not controlled)